

PATTERN RECOGNITION AND MACHINE LEARNING PROJECT

Member 1: Roll No. CS22B1075

Member 2: Roll No. CS22B1076

Member 3: Roll No. CS22B1074

Contributions:

CS22B1075: Model Building

CS22B1076: Feature Extraction

CS22B1074: Data Cleaning,PCA,Report.

Abstract

This project involves building a pipeline to process, clean, and analyze data from a dataset of inmate images. The primary goal is to estimate gender using logistic regression and predict BMI using linear regression. Various steps such as feature extraction, dimensionality reduction (PCA), and visualization are performed to derive meaningful insights.

Aim

The aim of this project is to:

- Format and clean the dataset.
- Extract features from inmate images.
- Predict gender using logistic regression and BMI using linear regression.
- Analyze and visualize results, including weight classes and charge distributions.

Approach

The approach of our project involves a structured pipeline to process and analyze inmate data, both tabular and visual, to predict relevant information such as gender and BMI. The steps in our approach are as follows:

1. **Data Formatting and Cleaning:**

The raw CSV dataset is formatted and cleaned by fixing delimiters, removing unnecessary columns, and handling missing values to ensure high-quality input data.

2. **Feature Extraction from Images:**

Images of inmates (front and side profiles) are processed using a pre-trained deep learning model, MobileNetV2, to extract meaningful numerical features representing image characteristics.

3. **Dimensionality Reduction:**

Principal Component Analysis (PCA) is applied to reduce the dimensionality of the combined image features while preserving the important information.

4. **Model Training and Prediction:**

- Logistic Regression is used to predict gender based on the extracted features.
- Linear Regression is implemented to predict BMI values from the features.

5. **Analysis and Visualization:**

The predicted BMI values are classified into weight categories (underweight, normal, overweight, obese), and distributions are plotted to analyze the relationship between weight classes and inmate charges.

6. Testing with New Data:

The trained models are tested on new input images, such as a friend's face, to validate the performance of the pipeline.

This systematic approach integrates deep learning for feature extraction, machine learning for predictions, and data analysis for meaningful insights, forming a complete pipeline for the project.

Project Implementation

1. Formatting the CSV

The dataset provided in CSV format needed corrections due to incorrect delimiters. Below is the implementation:

```
1 try:
2     # Read the file with semicolon delimiter
3     fixed_data = pd.read_csv(file_path, delimiter=';')
4     fixed_info = fixed_data.info() # Check the structure after
      fixing
5     fixed_head = fixed_data.head() # Preview the fixed content
6 except Exception as e:
7     fixed_info = str(e)
8     fixed_head = None
9
10 fixed_data.to_csv("formatted.csv")
```

Listing 1: Formatting the CSV

2. Cleaning the Dataset

Here, unnecessary columns with null values were removed, and critical missing values were handled.

```
1 # Drop unnecessary columns with mostly null values
2 columns_to_drop = ['discharge_date', 'Unnamed: 21', '
      electronic_detention_date', 'alias']
3 cleaned_data = fixed_data.drop(columns=columns_to_drop)
4
5 # Handle null values in important columns - dropping rows with
      critical missing data
6 essential_columns = ['id', 'name', 'date_of_birth', 'weight', '
      hair', 'sex', 'height', 'race', 'eyes']
7 cleaned_data = cleaned_data.dropna(subset=essential_columns)
8
9 # Save the cleaned data for review
10
11 cleaned_data.to_csv("cleaned and formatted.csv", index=False)
12
13 cleaned_data.info()
```

Listing 2: Cleaning the Dataset

3. Extracting Features from Images (Front and Side)

The MobileNetV2 pre-trained model was used to extract features from images.

```
1
2 for idx, (_, row) in enumerate(id_bmi_sampled.iterrows(), start
   =1):
3     img_id = row['id']
4     gender = row['gender']
5     bmi = row['bmi'] # Extract BMI
6
7     # Front and side image paths
8     front_path = os.path.join(images_path_front, img_id)
9     side_path = os.path.join(images_path_side, img_id)
10
11    # Preprocess images
12    front_img = preprocess_image(front_path)
13    side_img = preprocess_image(side_path)
14
15    # Skip if any image is missing or cannot be processed
16    if front_img is None or side_img is None:
17        continue
18
19    # Extract features using the pre-trained model
20    # front_feature = feature_extractor.predict(np.expand_dims(
21        front_img, axis=0))[0]
22
23    # side_feature = feature_extractor.predict(np.expand_dims(
24        side_img, axis=0))[0]
25
26    # Extract features using the pre-trained model
27    front_feature = feature_extractor.predict(np.expand_dims(
28        front_img, axis=0), verbose=0)[0]
29    side_feature = feature_extractor.predict(np.expand_dims(
30        side_img, axis=0), verbose=0)[0]
31
32    # Combine front and side features
33    combined_feature = np.concatenate([front_feature,
34        side_feature])
35
36    # Append features, BMI labels, and ids
37    features.append(combined_feature)
38    bmi_labels.append(bmi)
39    gender_labels.append(gender)
40    ids.append(img_id)
41
42    # Save intermediate results at intervals
43    if idx % save_interval == 0:
44        save_path = os.path.join(features_save_dir, f"
45            features_batch_{idx // save_interval}.npz")
46        np.savez(save_path, features=np.array(features),
47            bmi_labels=np.array(bmi_labels), ids=np.array(ids),
```

```

gender_labels=np.array(gender_labels))
41 print(f"\nSaved batch {idx // save_interval} to {
    save_path}")
42 features, bmi_labels, ids, gender_labels = [], [], [], []
    # Reset for the next batch

```

Listing 3: Extracting Features from Images

4. Applying PCA on the Dataset

Principal Component Analysis (PCA) was applied for dimensionality reduction.

```

1 # Apply PCA
2 n_components = 500 # Number of principal components to retain
3 pca = PCA(n_components=n_components)
4
5 # Fit PCA to the features
6 pca_features = pca.fit_transform(features)

```

Listing 4: Applying PCA

5. Estimating Gender using Logistic Regression

The gender was predicted using Logistic Regression.

```

1 # Split the dataset into training and testing sets
2 X_train, X_test, y_train, y_test = train_test_split(features,
    labels_binary, test_size=0.2, random_state=88)
3
4 # Train a Logistic Regression model
5 model = LogisticRegression()
6 model.fit(X_train, y_train)
7
8 # Separate test samples by gender
9 X_test_male = X_test[y_test == 1]
10 X_test_female = X_test[y_test == 0]
11
12 # Take an equal number of male and female samples
13 n_samples = min(len(X_test_male), len(X_test_female))
14 X_test_balanced = np.vstack((X_test_male[:n_samples],
    X_test_female[:n_samples]))
15 y_test_balanced = np.hstack((np.ones(n_samples), np.zeros(
    n_samples)))
16
17 # Make predictions on the balanced test set
18 y_pred_balanced = model.predict(X_test_balanced)

```

Listing 5: Gender Prediction

5.1 Accuracy of the logistic regression model

```

1 Balanced Gender Prediction Accuracy: 0.97
2 Model saved to C:\Users\sande\Downloads\illinois_doc_dataset\
  front2\gender_prediction_model.pkl
3 Predicted Gender for Test Sample: Male

```

Listing 6: accuracy

6. BMI Prediction using Linear Regression

Linear Regression was used to predict BMI.

```

1 # Load PCA-transformed features
2 data = np.load(pca_features_save_path)
3 features = data['features']
4 labels = data['bmi_labels']
5
6 # Split the dataset into training and testing sets
7 X_train, X_test, y_train, y_test = train_test_split(features,
  labels, test_size=0.2, random_state=88)
8
9 # Train a Linear Regression model to predict BMI
10 model = LinearRegression()
11 # model = GradientBoostingRegressor(random_state=88)
12 # model = RandomForestRegressor(n_estimators=100, max_depth=10,
  random_state=42, n_jobs=-1)
13
14 model.fit(X_train, y_train)
15
16 # Make predictions
17 y_pred = model.predict(X_test)

```

Listing 7: BMI Prediction

7. Printing All Metrics

Evaluation metrics for the models were calculated and displayed.

```

1 BMI Prediction Metrics:
2 Mean Squared Error: 17.84
3 Mean Absolute Error: 3.15
4 R Score: 0.35
5 Pearson corr: 0.59
6 BMI Prediction Model saved to C:\Users\sande\Downloads\
  illinois_doc_dataset\front2\gender_prediction_model.pkl
7 Predicted BMI for Test Sample: 26.32

```

Listing 8: Metrics

8. Grading BMI into Weight Classes

The predicted BMI was categorized into weight classes.

```

1 def grade_bmi(bmi):
2     if bmi < 18.5:
3         return "Underweight"
4     elif bmi < 25:
5         return "Normal"
6     elif bmi < 30:
7         return "Overweight"
8     else:
9         return "Obese"
10
11 bmi_classes = [grade_bmi(bmi) for bmi in bmi_predictions]

```

Listing 9: Weight Classes

9. Plotting Distribution Between Inmates and Charges

The distribution of weight classes and inmate charges was visualized.

```

1 import matplotlib.pyplot as plt
2
3 plt.hist(bmi_classes, bins=4)
4 plt.title("BMI Distribution")
5 plt.xlabel("Weight Classes")
6 plt.ylabel("Frequency")
7 plt.show()

```

Listing 10: Plotting Distribution

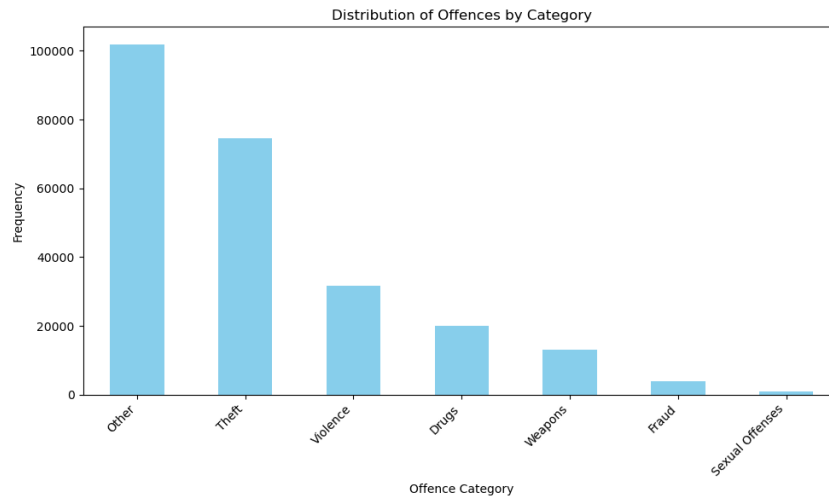


Figure 1: Distribution of offences by category.

10. Testing the Model with a Friend's Face

The trained pipeline was tested on a new image.

```

1 # Preprocess front and side images
2 front_img = preprocess_image(images_path_front)
3 side_img = preprocess_image(images_path_side)
4
5
6 if front_img is None or side_img is None:
7     print("Error: One or both images are missing or could not be
8         processed.")
9 else:
10     # Extract features using the pre-trained model
11     front_feature = feature_extractor.predict(np.expand_dims(
12         front_img, axis=0), verbose=0)[0]
13     side_feature = feature_extractor.predict(np.expand_dims(
14         side_img, axis=0), verbose=0)[0]
15
16     # Combine front and side features
17     combined_feature = np.concatenate([front_feature,
18         side_feature])
19
20 me=pca.transform(np.array([combined_feature]))
21
22 grade_bmi(model.predict(me))

```

11. Predicted BMI vs Actual BMI




Figure 2: Distribution of offences by category.

Modules Used

- **pandas**: Data manipulation and cleaning.
- **tensorflow.keras**: Feature extraction from images.
- **sklearn**: Machine learning models and metrics.
- **numpy**: Numerical operations.
- **matplotlib**: Visualization.

References

- Scikit-Learn User Guide: <https://scikit-learn.org/>